

# APL 405: Machine Learning for Mechanics

## Lecture 15: Convolutional Neural Network

by


Rajdip Nayek

Assistant Professor,  
Applied Mechanics Department,  
IIT Delhi

Instructor email: [rajdipn@am.iitd.ac.in](mailto:rajdipn@am.iitd.ac.in)


# Introduction

- We looked at **fully connected** neural networks which has each unit of previous layer is connected to all other units of the next layer



- **Drawbacks** of fully connected neural nets:
  - There are a **lot of connections**. Ex.  $p$  units in previous layer,  $q$  units in the next layer, then  $pq$  connections
  - If we are trying to classify an image, we flatten the 2D image into vectors, which **discards the spatial structure/information** of the image
- When dealing with images, the nearby pixels are typically related to each other, and we want to **exploit this neighbourhood (or local) information** to build more efficient neural networks

# From fully-connected layers to Convolution layers



- Suppose we want to train a network (with 1000 FC hidden units) that takes a  $200 \times 200$  colored (RGB) image as input

- What is the problem?

- **Too many parameters!** (Very complex, more chance of overfitting)

Input size =  $200 \times 200 \times 3 = 1,20,000$


Parameters =  $1,20,000 \times 1000 = 12 \times 10^7$

# Grayscale vs Colored images

Grayscale Image



height



width


- Grayscale images have a single channel (depth = 1)
- Colored images have more than one channel (depth > 1)
- Ex. RGB images has **3 channels**

Color Image



height

width



depth

# From fully-connected layers to Convolution layers

- Suppose we want to train a network (with 1000 hidden units) that takes a  $200 \times 200$  RGB image as input

- What is the problem?

- **Too many parameters!** (Very complex, more chance of overfitting)

Input size =  $200 \times 200 \times 3 = 1,20,000$

Parameters =  $1,20,000 \times 1000 = 12 \times 10^7$




- **Too translation sensitive** — Precise locations of objects in the image matter too much
  - If you translate the objects in the image to different locations, you may have to re-train a fully-connected MLP, else it may fail to classify the outputs correctly
- We can do much better with CNN for images



# From fully-connected layers to Convolution layers



- “Where’s Waldo?”  
In the game, Waldo shows up somewhere in some unlikely location. The reader’s goal is to locate him






# From fully-connected layers to Convolution layers



- “Where’s Waldo?”  
In the game, Waldo shows up somewhere in some unlikely location. The reader’s goal is to locate him
- We could sweep the image with a **Waldo detector** that could assign a score to each patch, indicating how likely the patch contains Waldo






# From fully-connected layers to Convolution layers



- “Where’s Waldo?”  
In the game, Waldo shows up somewhere in some unlikely location. The reader’s goal is to locate him
- We could sweep the image with a **Waldo detector** that could assign a score to each patch, indicating how likely the patch contains Waldo






# From fully-connected layers to Convolution layers



- “Where’s Waldo?”  
In the game, Waldo shows up somewhere in some unlikely location. The reader’s goal is to locate him
- We could sweep the image with a **Waldo detector** that could assign a score to each patch, indicating how likely the patch contains Waldo
- The patch with maximum score is where Waldo should be located
- As this local patch sweeps the entire image, it does not matter where Waldo is located



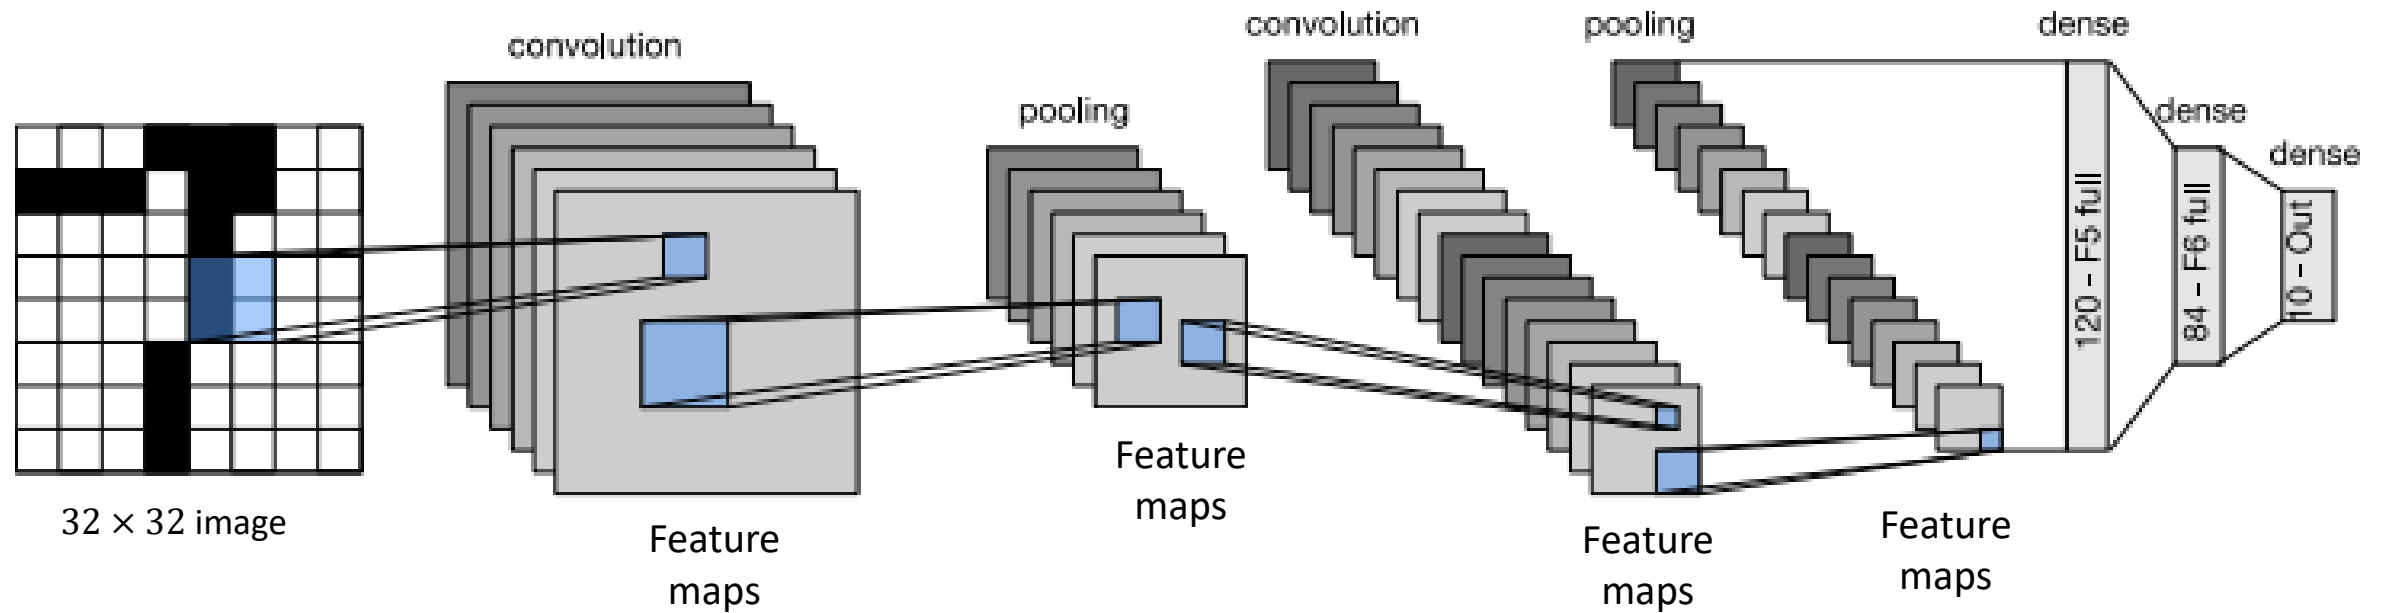
# From fully-connected layers to Convolution layers



- “Where’s Waldo?”  
In the game, Waldo shows up somewhere in some unlikely location. The reader’s goal is to locate him
- We could sweep the image with a **Waldo detector** that could assign a score to each patch, indicating how likely the patch contains Waldo
- CNNs systematize this idea of **translation invariance** and **localised feature detection**, via convolutions and max pooling, with much less parameters



# From fully-connected layers to Convolution layers



- CNNs systematize this idea of **translation invariance** and **localised feature detection**, via **convolutions** and **max pooling**, with much less parameters
- CNNs uses multiple **kernels** (“Waldo detectors”) that detects different features

What is convolution?



# 1D Convolution

- Convolution of two scalar-valued functions  $w(x)$  and  $g(x)$  is defined as:  $s(x) = (w * g)(x) = \int w(x - a) g(a) da$

- Whenever we have discrete objects (arrays), the integral turns into a sum:  $s[i] = \sum_a w[i - a] g[a]$

$w[0]$	$w[1]$	$w[2]$	$w[3]$		$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$		
0.1	0.2	0.3	0.4	*	1	2	3	4	5	6	=	?

- The array  $g$  is the input
- The array  $w$  is called the **filter (or kernel)**

# 1D Convolution

- Convolution of two scalar-valued functions  $w(x)$  and  $g(x)$  is defined as:  $s(x) = (w * g)(x) = \int w(x - a) g(a) da$

- Whenever we have discrete objects (arrays), the integral turns into a sum:  $s[i] = \sum_a w[i - a] g[a]$

$w[0]$	$w[1]$	$w[2]$	$w[3]$
0.1	0.2	0.3	0.4

\*

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

=
?

- The array  $g$  is the input
- The array  $w$  is called the **filter (or kernel)**
- Flip-and-filter**

- Slide the filter over the input and compute windowed dot product

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

$w[3]$	$w[2]$	$w[1]$	$w[0]$
0.4	0.3	0.2	0.1

(flipped)

$$s[3] = w[3]g[0] + w[2]g[1] + w[1]g[2] + w[0]g[3]$$

$s[3]$	$s[4]$	$s[5]$
2		



# Convolution in 1D

- Convolution of two scalar-valued functions  $w(x)$  and  $g(x)$  is defined as:  $s(x) = (w * g)(x) = \int w(x - a) g(a) da$

- Whenever we have discrete objects (arrays), the integral turns into a sum:  $s[i] = \sum_a w[i - a] g[a]$

$w[0]$	$w[1]$	$w[2]$	$w[3]$
0.1	0.2	0.3	0.4

\*

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

=
?

- The array  $g$  is the input
- The array  $w$  is called the **filter (or kernel)**
- Flip-and-filter**

- Slide the filter over the input and compute windowed dot product

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

$w[3]$	$w[2]$	$w[1]$	$w[0]$
0.4	0.3	0.2	0.1

$$s[4] = w[3]g[1] + w[2]g[2] + w[1]g[3] + w[0]g[4]$$

$s[3]$	$s[4]$	$s[5]$
2	4	

# 1D Convolution

- Convolution of two scalar-valued functions  $w(x)$  and  $g(x)$  is defined as:  $s(x) = (w * g)(x) = \int w(x - a) g(a) da$

- Whenever we have discrete objects (arrays), the integral turns into a sum:  $s[i] = \sum_a w[i - a] g[a]$

$w[0]$	$w[1]$	$w[2]$	$w[3]$
0.1	0.2	0.3	0.4

\*

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

=
?

- The array  $g$  is the input
- The array  $w$  is called the **filter (or kernel)**

- Flip-and-filter**

- Slide the filter over the input and compute windowed dot product

- Here the input (and the kernel) is 1D

$g[0]$	$g[1]$	$g[2]$	$g[3]$	$g[4]$	$g[5]$
1	2	3	4	5	6

$w[3]$	$w[2]$	$w[1]$	$w[0]$
0.4	0.3	0.2	0.1

$$s[5] = w[3]g[2] + w[2]g[3] + w[1]g[4] + w[0]g[5]$$

$s[3]$	$s[4]$	$s[5]$
2	4	5

# 1D Convolution to 2D Convolution


- Convolution is more like doing a **flipped cross-correlation** operation
- The **filters (or kernels)** will resemble the weights in CNN (as we will see soon)
- Most machine learning libraries just implement a moving window cross-correlation (and *ignore flipping*) since it does not matter much whether you learn a flipped set of weights or unflipped set of weights
- How does convolution (think more like cross-correlation) look in 2D?
- Let's now consider 2D grayscale images (has depth of 1) and 2D kernels

Input		Kernel		Output																	
<table border="1"><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td><td>5</td></tr><tr><td>6</td><td>7</td><td>8</td></tr></table>	0	1	2	3	4	5	6	7	8	*	<table border="1"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	0	1	2	3	=	<table border="1"><tr><td></td><td></td></tr><tr><td></td><td></td></tr></table>				
0	1	2																			
3	4	5																			
6	7	8																			
0	1																				
2	3																				
$3 \times 3$		$2 \times 2$		$2 \times 2$																	



# 1D Convolution to 2D Convolution

- Convolution is more like doing a **flipped cross-correlation** operation
- The **filters (or kernels)** will resemble the weights in CNN (as we will see soon)
- Most machine learning libraries just implement a moving window cross-correlation (and *ignore flipping*) since it does not matter much whether you learn a flipped set of weights or unflipped set of weights
- How does convolution (think more like cross-correlation) look in 2D?
- Let's now consider 2D grayscale images (has depth of 1) and 2D kernels



$$0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 = 19$$

# 2D Convolution

- Convolution is more like a moving window **flipped cross-correlation** operation
- The **filters (or kernels)** will resemble the weights in CNN (as we will see soon)
- Most machine learning libraries just implement a moving window cross-correlation (and *ignore flipping*) since it does not matter much whether you learn a flipped set of weights or unflipped set of weights
- How does convolution (think cross-correlation from now on) look in 2D?
- Let's now consider 2D grayscale images (has depth of 1) and 2D kernels

Input				Kernel			Output	
0	1	2	*	0	1	=	19	25
3	4	5		2	3			
6	7	8						

$$0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 = 19$$

$$1 \times 0 + 2 \times 1 + 4 \times 2 + 5 \times 3 = 25$$

# 2D Convolution

- Convolution is more like a moving window **flipped cross-correlation** operation
- The **filters (or kernels)** will resemble the weights in CNN (as we will see soon)
- Most machine learning libraries just implement a moving window cross-correlation (and *ignore flipping*) since it does not matter much whether you learn a flipped set of weights or unflipped set of weights
- How does convolution (think cross-correlation from now on) look in 2D?
- Let's now consider 2D grayscale images (has depth of 1) and 2D kernels

Input				Kernel			Output	
0	1	2		0	1		19	25
3	4	5	*	2	3	=	37	
6	7	8						

$$\begin{aligned} 0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 &= 19 \\ 1 \times 0 + 2 \times 1 + 4 \times 2 + 5 \times 3 &= 25 \\ 3 \times 0 + 4 \times 1 + 6 \times 2 + 7 \times 3 &= 37 \end{aligned}$$



# 2D Convolution

- Convolution is more like a moving window **flipped cross-correlation** operation
- The **filters (or kernels)** will resemble the weights in CNN (as we will see soon)
- Most machine learning libraries just implement a moving window cross-correlation (and *ignore flipping*) since it does not matter much whether you learn a flipped set of weights or unflipped set of weights
- How does convolution (think cross-correlation from now on) look in 2D?
- Let's now consider 2D grayscale images (has depth of 1) and 2D kernels

Input				Kernel			Output		
0	1	2	*	0	1	=	19	25	$0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 = 19$ $1 \times 0 + 2 \times 1 + 4 \times 2 + 5 \times 3 = 25$ $3 \times 0 + 4 \times 1 + 6 \times 2 + 7 \times 3 = 37$ $4 \times 0 + 5 \times 1 + 7 \times 2 + 8 \times 3 = 43$
3	4	5		2	3		37	43	
6	7	8							

# 2D Convolution

- Despite the simplicity of the operation, convolution can do some pretty interesting things
- Let's look at some examples of convolutions with grayscale images



\*

1	1	1
1	1	1
1	1	1

=



**Blur kernel:** Takes an average of all the neighbouring pixels

# 2D Convolution

- Despite the simplicity of the operation, convolution can do some pretty interesting things
- Let's look at some examples of convolutions with grayscale images



\*

0	-1	0
-1	5	-1
0	-1	0

=



**Sharpen kernel:** Emphasizes differences in adjacent pixel values



# 2D Convolution

- Despite the simplicity of the operation, convolution can do some pretty interesting things
- Let's look at some examples of convolutions with grayscale images



\*

-1	0	1
-2	0	2
-1	0	1



=



**Vertical edge detector:** Finds edges from darker to brighter intensities

# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution

-1	0	1
-2	0	2
-1	0	1





- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output



# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution

-1	0	1
-2	0	2
-1	0	1





- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output



# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution

-1	0	1
-2	0	2
-1	0	1





- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output



# 2D Convolution



-1	0	1
-2	0	2
-1	0	1



- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output

# 2D Convolution

-1	0	1
-2	0	2
-1	0	1



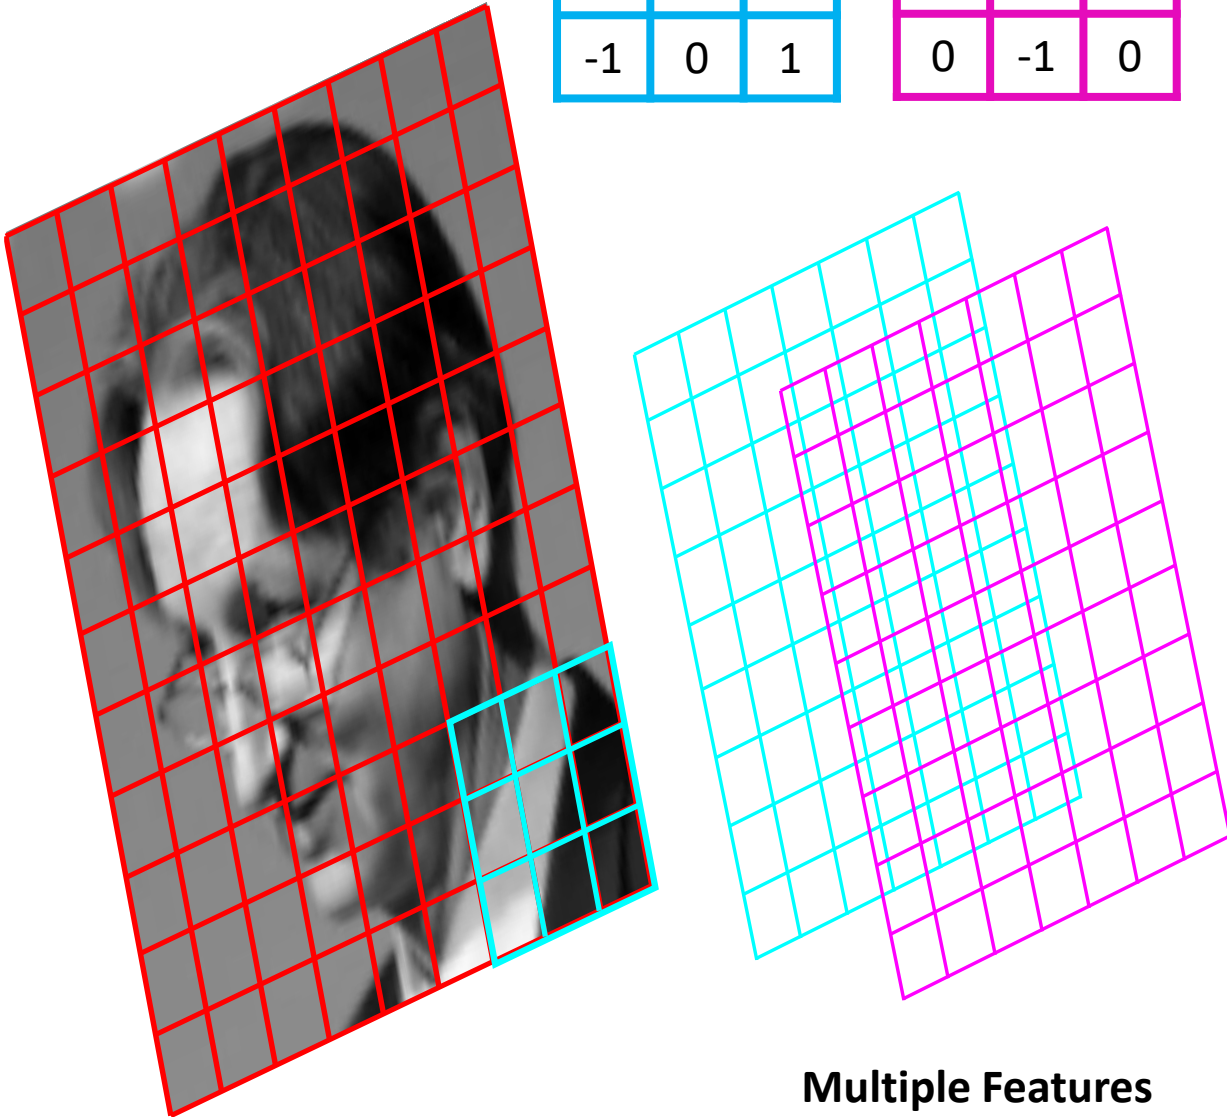
Feature map

- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output
- The resulting output is called a **feature map**

# 2D Convolution

-1	0	1
-2	0	2
-1	0	1

0	-1	0
-1	5	-1
0	-1	0



Multiple Features

- We just slide the kernel over the input image
- Each time we slide the kernel we get one value in the output
- The resulting output is called a **feature map**
- We can use **multiple filters** to get multiple feature maps
- How convolutions will happen for colored images?




# What would happen in case of colored images?

Grayscale Image



height




width

Color Image




height

width




- Grayscale images have a single channel (depth = 1)
- Colored images have more than one channel (depth > 1)
- Ex. RGB images has **3 channels**
- How does convolution happen in 3D case?
- Well the kernel or filter will be 3D too (i.e. will have same number of channels as the input)

# What would happen in case of colored images?



Input

- Grayscale images have a single channel (depth = 1)
- Colored images have more than one channel (depth > 1)
- Ex. RGB images has **3 channels**
- How does convolution happen in 3D case?
- Well the kernel or filter will be 3D too (i.e. will have same number of channels as the input)



A single 3D kernel

- Let's see how the 3D convolutions happen

		-6	-2	3	0
	-3	1	2	3	4
-2	1	5	4	2	8
2	5	6	7	8	
6	7	8	9		

RGB Input (3 channels)

\*

			-1	1	2
			3	2	0
	0		6	0	1
0	3	0	5		
3	2	-1	8		
6	7	8			

A single 3D kernel

-6	-2	3	0
1	2	3	4
5	4	2	8

\*

-1	1	2
3	2	0
6	0	1

+

-2	-1	0	1
2	3	4	5
6	7	8	9

\*

0	1	2
3	4	5
6	7	8

+

-3	-2	-1	0
1	2	3	4
5	6	7	8

\*

0	-4	0
3	0	5
2	-1	8

-2	-1	0	1
2	3	4	5
6	7	8	9

+

-3	-2	-1	0
1	2	3	4
5	6	7	8

+

-6	-2	3	0
1	2	3	4
5	4	2	8

\*

0	1	2
3	4	5
6	7	8

\*

0	-4	0
3	0	5
2	-1	8

\*

-1	1	2
3	2	0
6	0	1

=

323	
-----	--

$$\begin{aligned}
 &(-2)(0) + (-1)(1) + (0)(2) \\
 &+ (2)(3) + (3)(4) + (4)(5) \\
 &+ (6)(6) + (7)(7) + (8)(8)
 \end{aligned}$$

188

+

$$\begin{aligned}
 &(-3)(0) + (-2)(-4) + (-1)(0) \\
 &+ (1)(3) + (2)(0) + (3)(5) \\
 &+ (5)(2) + (6)(-1) + (7)(8)
 \end{aligned}$$

86

+

$$\begin{aligned}
 &(-6)(-1) + (-2)(1) + (3)(2) \\
 &+ (1)(3) + (2)(2) + (3)(0) \\
 &+ (5)(6) + (4)(0) + (2)(1)
 \end{aligned}$$

49



-2	-1	0	1
2	3	4	5
6	7	8	9

+

-3	-2	-1	0
1	2	3	4
5	6	7	8

+

-6	-2	3	0
1	2	3	4
5	4	2	8

\*

0	1	2
3	4	5
6	7	8

\*

0	-4	0
3	0	5
2	-1	8

\*

-1	1	2
3	2	0
6	0	1

=

323	370
-----	-----

$(-1)(0) + (0)(1) + (1)(2)$   
 $+ (3)(3) + (4)(4) + (5)(5)$   
 $+ (7)(6) + (8)(7) + (9)(8)$

**222**

+

$(-2)(0) + (-1)(-4) + (0)(0)$   
 $+ (2)(3) + (3)(0) + (4)(5)$   
 $+ (6)(2) + (7)(-1) + (8)(8)$


**99**

+

$(-2)(-1) + (3)(1) + (0)(2)$   
 $+ (2)(3) + (3)(2) + (4)(0)$   
 $+ (4)(6) + (2)(0) + (8)(1)$

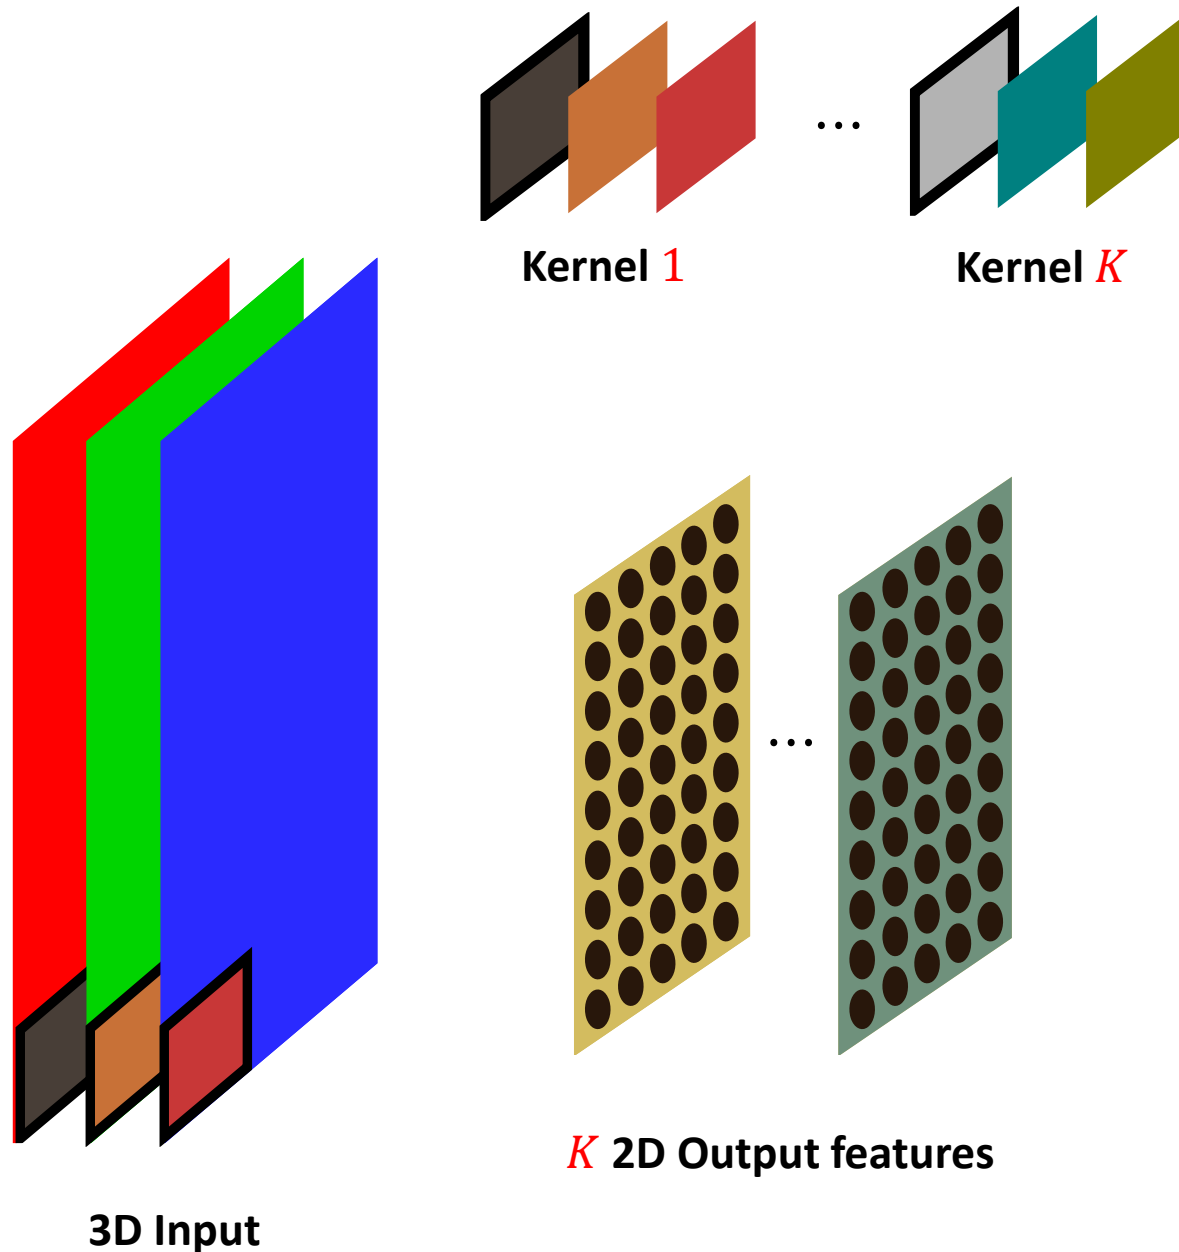
**49**

# What would happen in case of colored images?



- Grayscale images have a single channel (depth = 1)
- Colored images have more than one channel (depth > 1)
- Ex. RGB images has **3 channels**
- **How does convolution happen in 3D case?**
- Well the kernel or filter will be 3D too (i.e. will have same number of channels as the input)
- The kernel moves along the width and height (and not along the depth)
- **Therefore, the feature output is 2D!**

# What would happen in case of colored images?



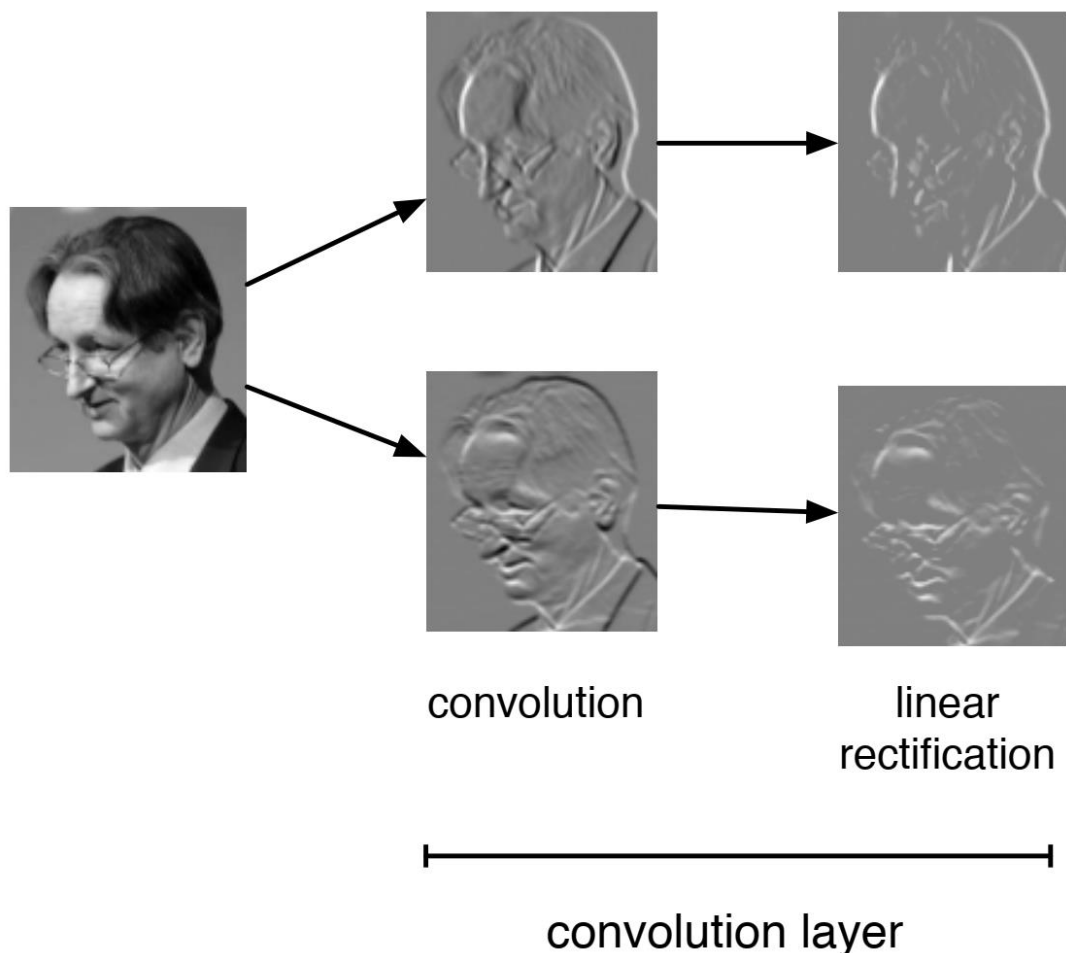
- Grayscale images have a single channel (depth = 1)
- Colored images have more than one channel (depth > 1)
- Ex. RGB images has **3 channels**
- **How does convolution happen in 3D case?**
- Well the kernel or filter will be 3D too (i.e. will have same number of channels as the input)
- The kernel moves along the width and height (and not along the depth)
- **Therefore, the feature output is 2D!**
- Once again, if we apply **multiple 3D filters**, we will get multiple 2D output features

# Convolution followed by linear rectification

It is common to apply a ReLU nonlinear activation on the output feature following convolution:  $y = \max(z, 0)$

Why might we do this?


- Convolution is a **linear** operation. Passing the linear output through nonlinear activation leads to more powerful features
- While pooling the results, two edges in opposite directions shouldn't cancel
- It has been reported that nonlinear activations (like ReLU or ELU) when used after convolutional layers given better performance






What are the relations between input sizes, kernel sizes, and output sizes?

# Relation between input sizes, kernel sizes, and output sizes




- So far we have not said anything explicit about the dimensions of the
  - Inputs
  - Kernels
  - Outputs
  - the relations between them
- We will see how they are related but before that we will discuss **zero-padding** and **stride**


# Zero padding



- Note that we can't place the kernel centred at the corners or at boundaries of our image
- Thus any interesting information on the boundaries of the original image is lost



# Zero padding



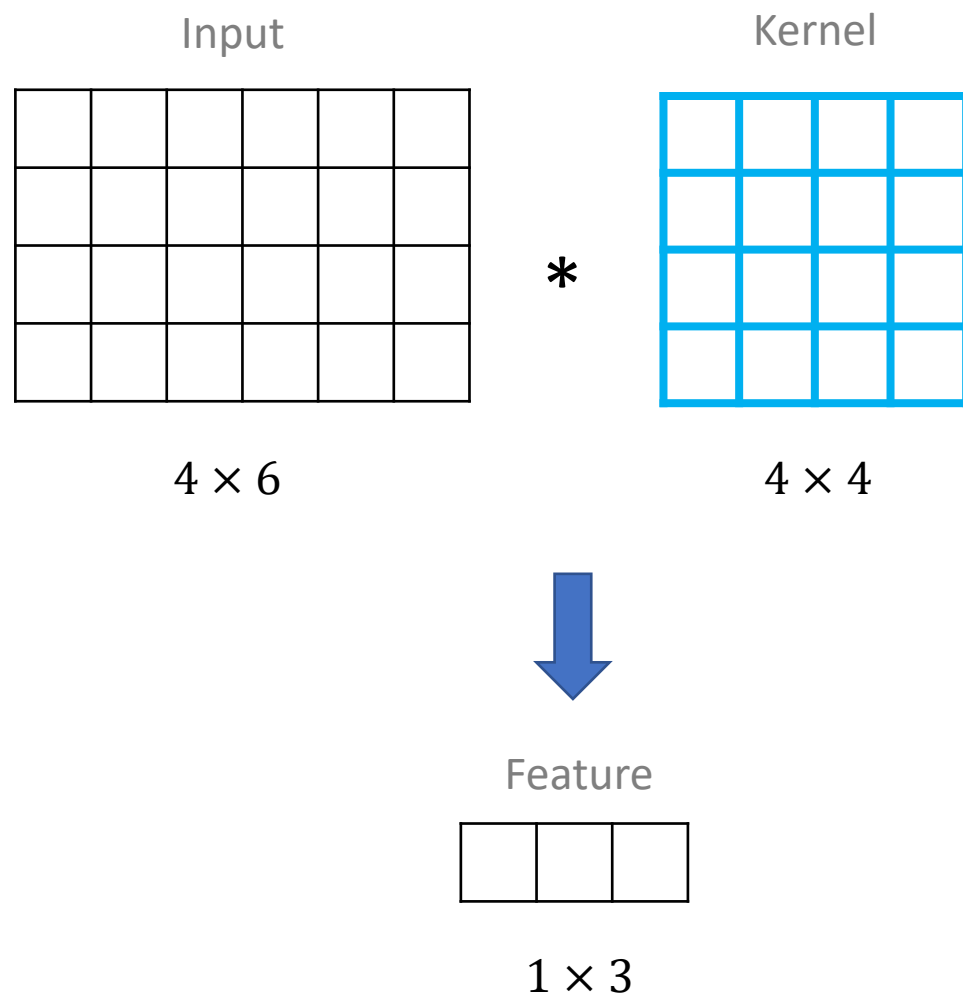
- Note that we can't place the kernel centred at the corners or at boundaries of our image
- Thus any interesting information on the boundaries of the original image is lost
- This loss of information results in an output feature size smaller than the input image
- If input size is  $n_h \times n_w$ , kernel size is  $k_h \times k_w$ , then output feature size  $f_h \times f_w$  is related as follows:

$$f_h = n_h - k_h + 1$$

$$f_w = n_w - k_w + 1$$



# Zero padding



- Note that we can't place the kernel centred at the corners or at boundaries of our image
- Thus any interesting information on the boundaries of the original image is lost
- This loss of information results in an output feature size smaller than the input image
- If input size is  $n_h \times n_w$ , kernel size is  $k_h \times k_w$ , then output feature size  $f_h \times f_w$  is related as follows:
$$f_h = n_h - k_h + 1$$
$$f_w = n_w - k_w + 1$$
- As the size of the kernel increases, this output size reduces even more

# Zero padding

Input


$4 \times 6$

Kernel


$3 \times 3$

Zero-padded Input

0	0	0	0	0	0	0	0
0							0
0							0
0							0
0							0
0	0	0	0	0	0	0	0

$6 \times 8$

- One straightforward solution to this problem is **to add zero pixels around the boundary of our input image**, thus increasing the effective size of the image
- This means we **pad zeros** of width  $p_w$  on left and right, and **pad zeros** of height  $p_h$  on top and bottom
- The output feature shape will be  $f_h \times f_w$ :

$$f_h = n_h - k_h + 2p_w + 1$$

$$f_w = n_w - k_w + 2p_h + 1$$

# Zero padding

Zero-padded Input

0	0	0	0	0	0	0	0
0							0
0							0
0							0
0							0
0	0	0	0	0	0	0	0

$6 \times 8$

Kernel


\*

$3 \times 3$

Output feature


$4 \times 6$

- One straightforward solution to this problem is **to add zero pixels around the boundary of our input image**, thus increasing the effective size of the image

- This means we pad zeros of width  $p_w$  on left and right, and pad zeros of height  $p_h$  on top and bottom

- The output feature shape will be  $f_h \times f_w$ :


$$f_h = n_h - k_h + 2p_w + 1$$

$$f_w = n_w - k_w + 2p_h + 1$$

- Usual values:  $p_h = \frac{k_h-1}{2}$  and  $p_w = \frac{k_w-1}{2}$

- It becomes easy to work with odd-sized kernels like 3x3, 5x5, 7x7 as zero-padding will give an integer number, else ceil the value

# Stride



- When computing the cross-correlation, we typically move our kernel by one interval along right and/or downwards
- **Stride** defines the intervals at which the kernel is applied
- By default, we slide one element at a time
- However, sometimes, **either for computational efficiency** or because **we wish to downsample (reduce size of output)**, we move our window more than one element at a time, skipping the intermediate locations
- For such cases, the stride would be greater than 1



# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . Let's consider an example with  $S = 2$

Input

0	0	0	0	0	0	0	0
0	1	2	3	4	4	5	0
0	6	0	7	8	9	0	0
0	4	1	7	2	9	4	0
0	9	4	9	1	4	2	0
0	0	0	0	0	0	0	0

$4 \times 6$

\*

Kernel

0	1	2
3	4	5
6	7	8

$3 \times 3$

$$\left\lfloor \frac{4 - 3 + 2}{2} + 1 \right\rfloor$$

$\lfloor 2.5 \rfloor$

$= 2$

Output feature


$2 \times 3$

$$\left\lfloor \frac{6 - 3 + 2}{2} + 1 \right\rfloor$$

$\lfloor 3.5 \rfloor$


$3$

# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained




# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained




# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained






# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained




# Stride

- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained




# Stride


- When the stride along the height is  $s_h$  and the stride along the width is  $s_w$ , the output shape  $f_h \times f_w$  is

$$\left\lfloor \frac{n_h - k_h + 2p_h}{s_h} + 1 \right\rfloor \times \left\lfloor \frac{n_w - k_w + 2p_w}{s_w} + 1 \right\rfloor$$

- Typically, equal strides are taken,  $s_h = s_w = S$ . For stride  $S = 2$ , the following output is obtained




# Depth of the output layer



- Finally, let's come to the depth  $f_d$  of the output feature layer
- If we have multi-channel inputs, depth will be  $n_d > 1$
- Each 3D kernel will give us one 2D output feature
- $K$  kernels will give us  $K$  such 2D output features
- We can think of the resulting output feature as  $f_h \times f_w \times f_d$  volume
- Thus,  $f_d = K$

# Example for determining sizes

- 6 kernels
- Stride,  $S = 1$
- Padding,  $P = 0$



- Output layer dimensions

$$f_h = \left\lfloor \frac{n_h - k_h + 2P}{S} + 1 \right\rfloor$$


$$f_w = \left\lfloor \frac{n_w - k_w + 2P}{S} + 1 \right\rfloor$$

$$f_d = K \text{ (number of kernels)}$$



# Example for determining sizes

- 96 kernels
- Stride,  $S = 4$
- Padding,  $P = 0$



- Output layer dimensions

$$f_h = \left\lfloor \frac{n_h - k_h + 2P}{S} + 1 \right\rfloor$$

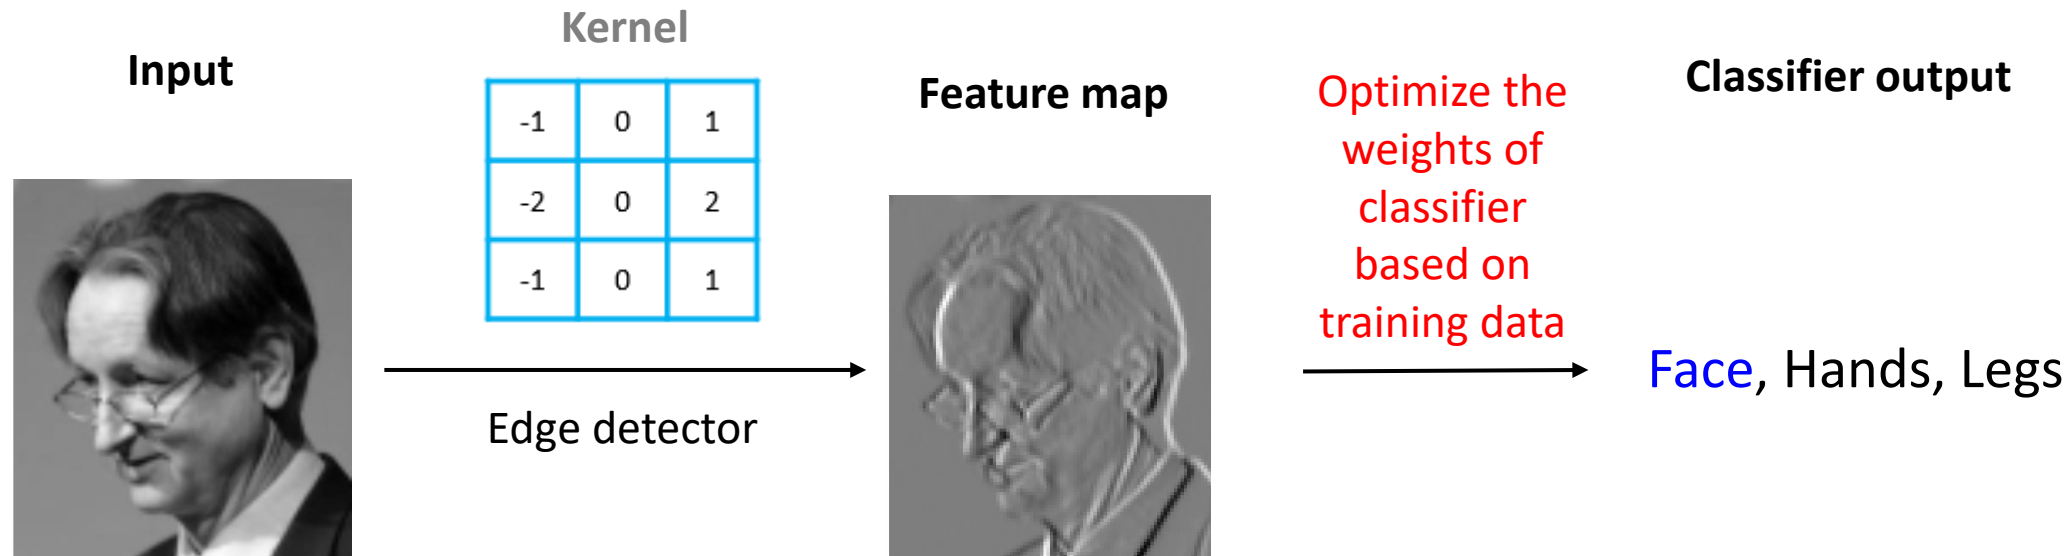
$$f_w = \left\lfloor \frac{n_w - k_w + 2P}{S} + 1 \right\rfloor$$

$$f_d = K \text{ (number of kernels)}$$

$$\begin{aligned} f_w &= \left\lfloor \frac{227 - 11}{4} + 1 \right\rfloor \\ &= \left\lfloor \frac{216}{4} + 1 \right\rfloor \\ &= \lfloor 54 + 1 \rfloor \\ &= 55 \end{aligned}$$


- What is the connection between this operation (convolution) and neural networks?
- We will try to understand this by considering the task of image classification

# Output features for image classification





- Instead of using handcrafted kernels such as edge detectors **can we learn (or optimize) meaningful kernels/filters in addition to learning the weights of the classifier?**
- Even better, if we can learn **multiple** meaningful kernels/filters in addition to the weights of the classifier
- In CNN, we treat these kernels as parameters and learn them in addition to the weights of the classifier (using back propagation) in CNN
- **But how is CNN different than fully-connected feedforward neural networks?**

# Full connected NN vs CNN



- This is what a regular fully-connected feed-forward neural network looks like
- It is **dense**, there are many connections
- For example, all the 16 input neurons are contributing to the computation of  $h_{11} = h_1^{(1)}$
- Contrast this to what happens in the case of convolution



# Full connected NN vs CNN



- Only a few local neurons participate in the computation of  $h_{11} = h_1^{(1)}$
- For example, only pixels 1, 2, 5, 6 contribute to  $h_1^{(1)}$





# Full connected NN vs CNN




- Only a few local neurons participate in the computation of  $h_{11} = h_1^{(1)}$
- For example, only pixels 1, 2, 5, 6 contribute to  $h_1^{(1)}$
- For example, only pixels 3, 4, 7, 8 contribute to  $h_{12} = h_2^{(1)}$

# Full connected NN vs CNN





- Only a few local neurons participate in the computation of  $h_{11} = h_1^{(1)}$
- For example, only pixels 1, 2, 5, 6 contribute to  $h_1^{(1)}$
- The connections are much sparser
- This **sparse connectivity** reduces the number of parameters in the model
- **But is sparse connectivity good?** Aren't we losing information (by losing interactions between some input pixels) ?

# Full connected NN vs CNN




- But is sparse connectivity good? Aren't we losing information (by losing interactions between some input pixels) ?
- It turns out we are not losing information/interactions
- The two highlighted neurons ( $x_1$  and  $x_5$ ) do not interact in *hidden layer 1*
- But they indirectly contribute to the computation of  $g_3$  and hence interact indirectly

# Full connected NN vs CNN



- Another characteristic of CNNs is **weight sharing**
- Imagine if we use an edge detection kernel
- Then the same kernel is passed over the all locations of the image to produce  $h_1^{(1)}, h_2^{(1)}, h_3^{(1)}, h_4^{(1)}, \dots$
- Since the kernel weights remain same as we sweep across all locations of the image, it is as if we share the weights across all locations of the image

# Full connected NN vs CNN

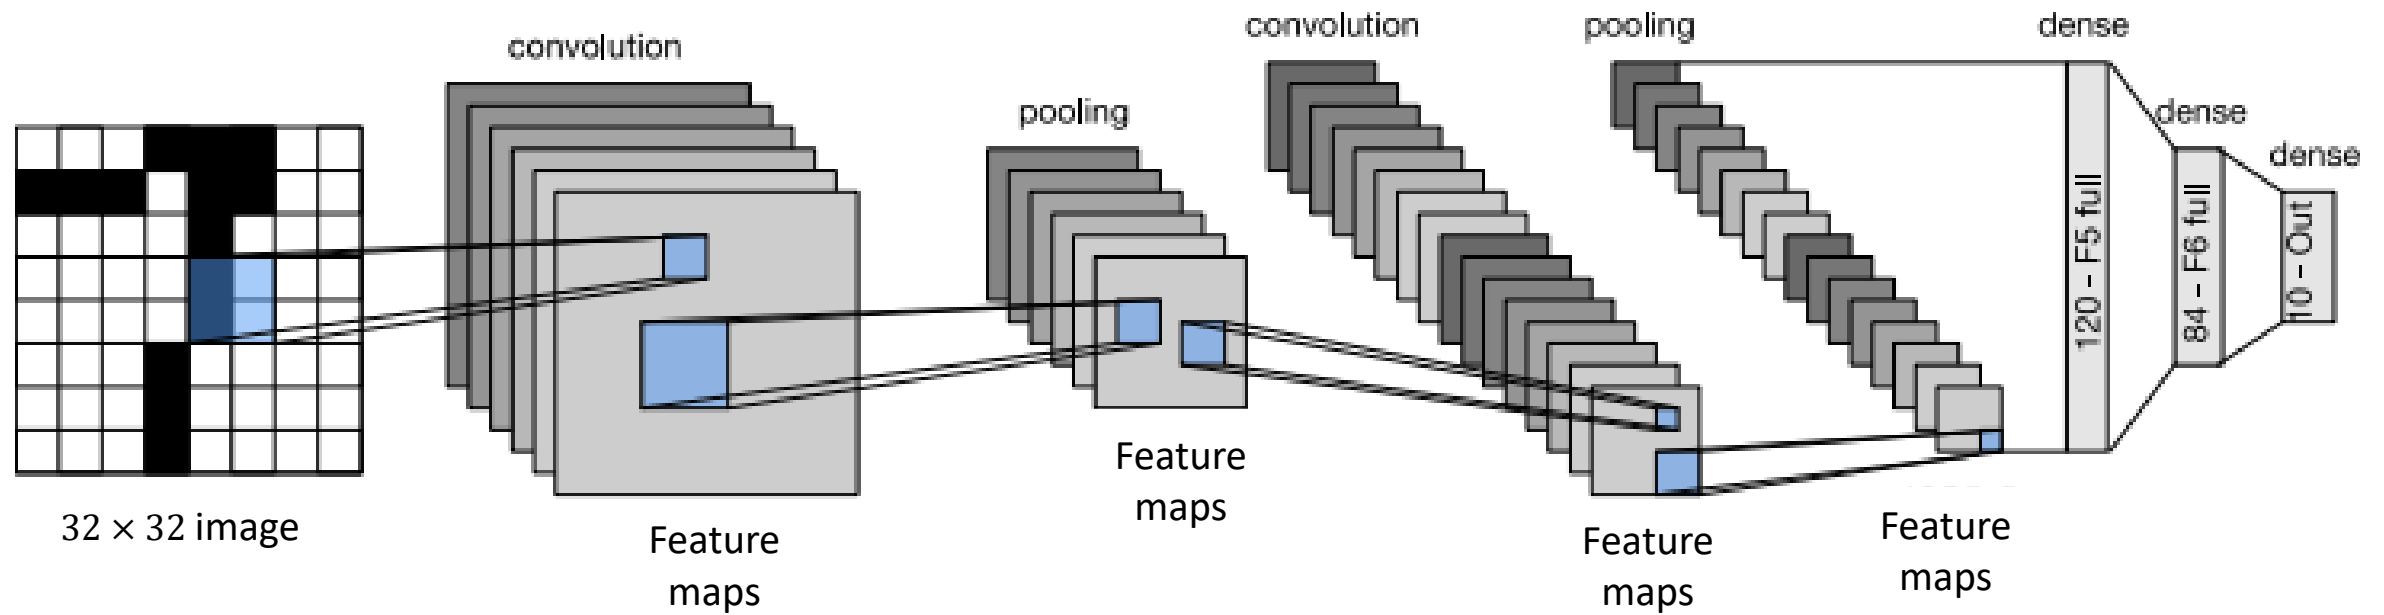


- Another characteristic of CNNs is **weight sharing**
- Imagine if we use an edge detection kernel
- Then the same kernel is passed over the all locations of the image to produce  $h_1^{(1)}, h_2^{(1)}, h_3^{(1)}, h_4^{(1)}, \dots$
- Since the kernel weights remain same as we sweep across all locations of the image, it is as if we share the weights across all locations of the image
- Note, we can have many such kernels and each kernel will be shared by all locations in the image

- So far we have talked a lot on convolution layers
- Saw how kernels are convolved with inputs to produce features
- Understood that kernels are to be learned (or optimized), not manually set
- Let's look at CNN for a moment

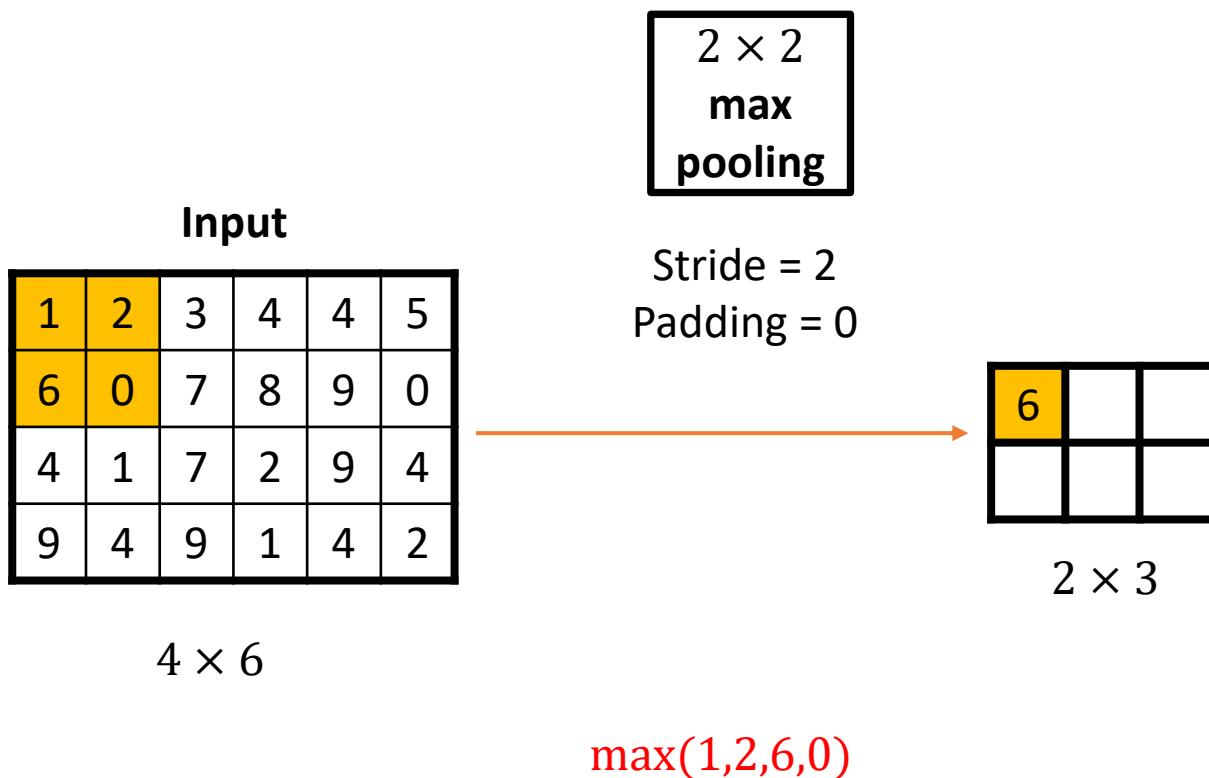


# Convolutional neural networks



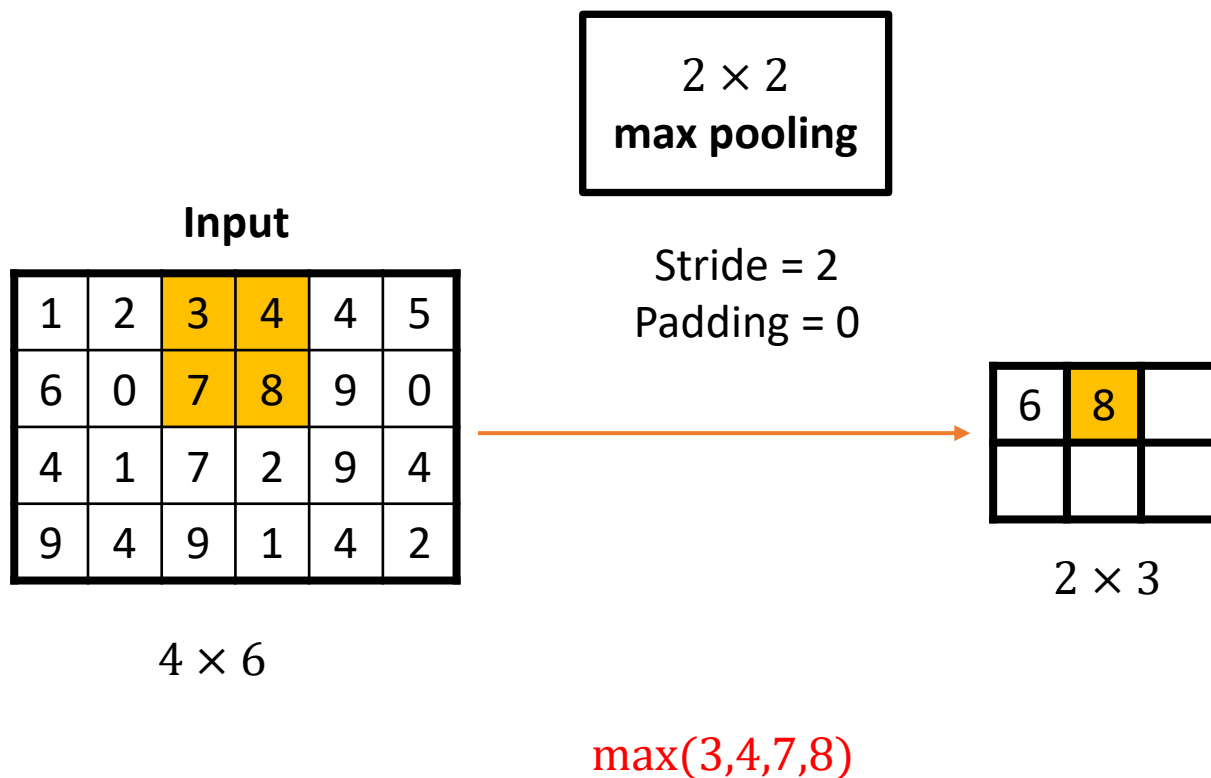
- So a CNN has alternate convolution and pooling layers
- What does a pooling layer do?

# Pooling



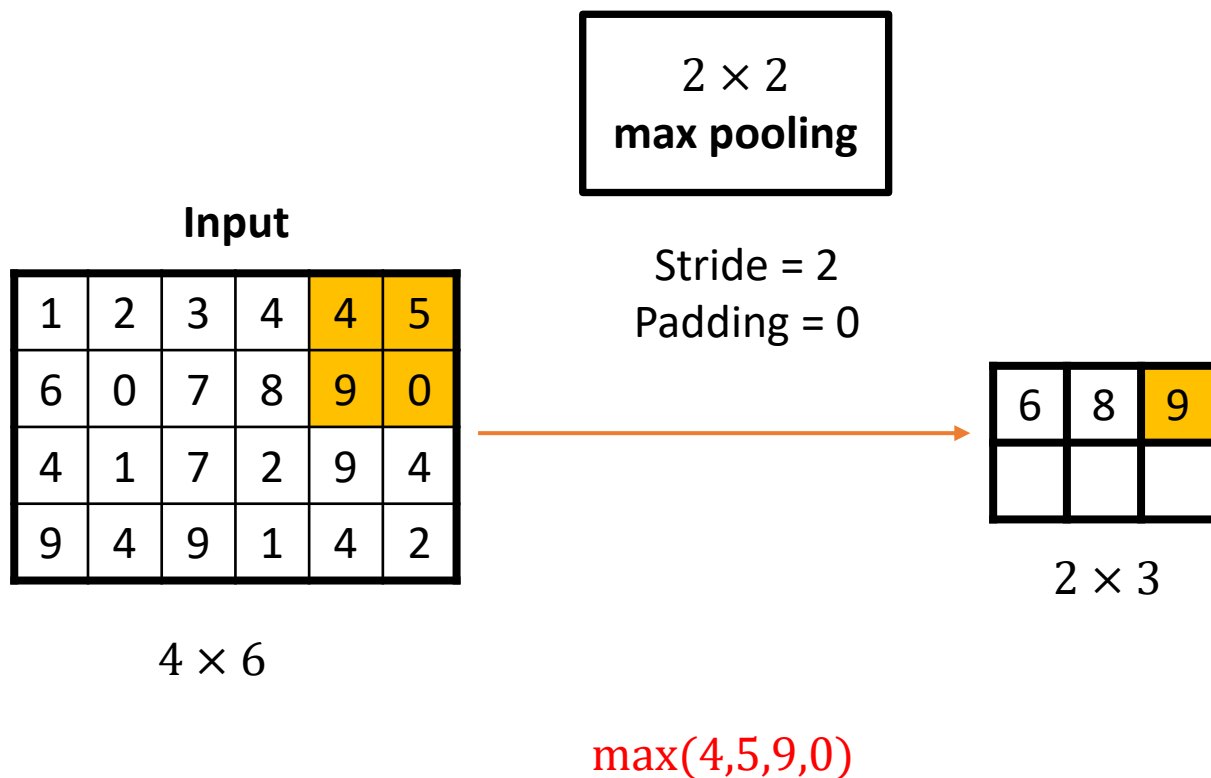
- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, **the pooling layer contains no parameters** (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**

# Pooling



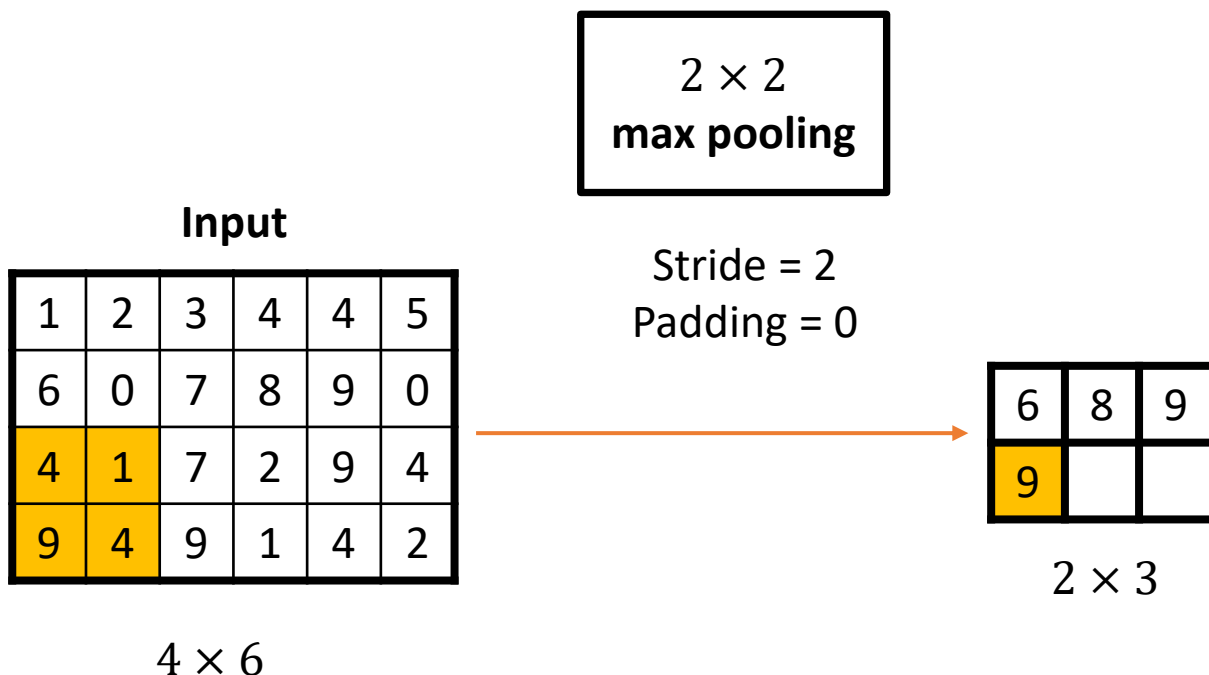
- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, the pooling layer contains no parameters (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**

# Pooling



- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, the pooling layer contains no parameters (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**

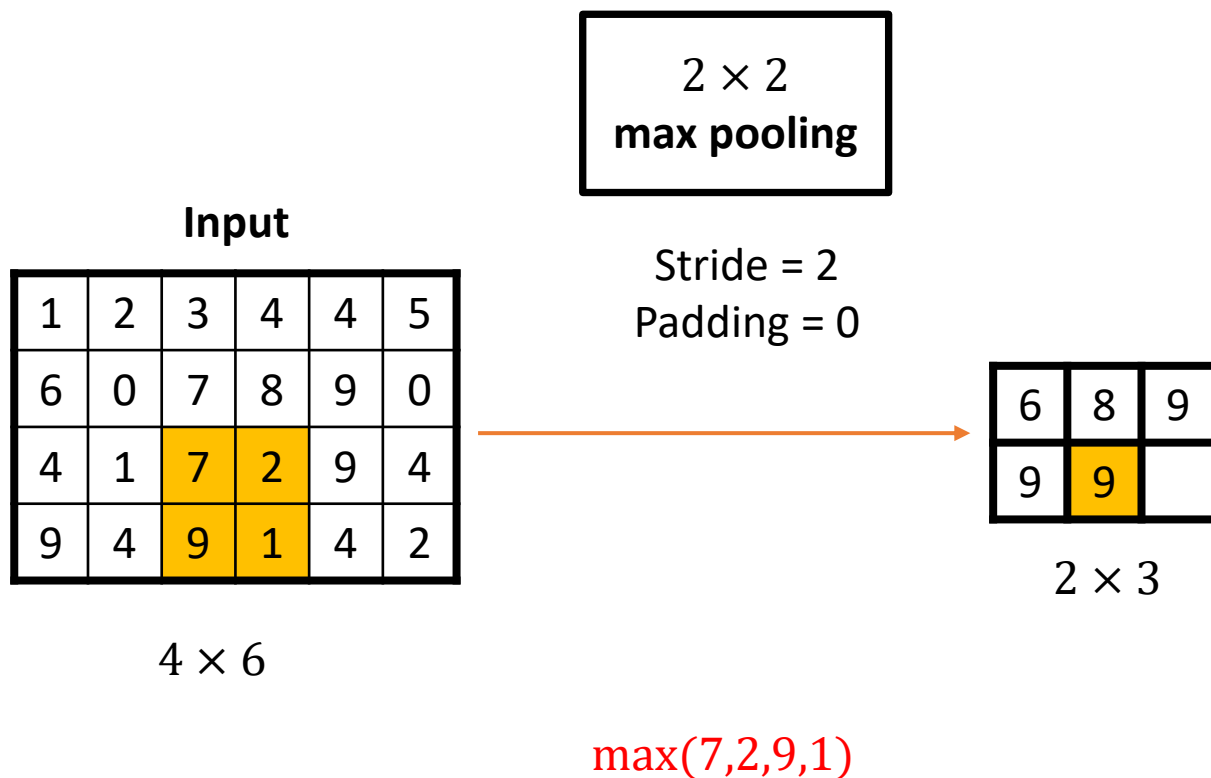
# Pooling



$\max(4, 1, 9, 4)$

- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, the pooling layer contains no parameters (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**

# Pooling



- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, the pooling layer contains no parameters (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**



# Pooling

Input

1	2	3	4	4	5
6	0	7	8	9	0
4	1	7	2	9	4
9	4	9	1	4	2

$4 \times 6$

$2 \times 2$   
**max pooling**

Stride = 2  
Padding = 0

6	8	9
9	9	9

$2 \times 3$


$\max(9, 4, 4, 2)$

- We want to gradually reduce the spatial resolution of our hidden representations while aggregating meaningful features
- Pooling helps in reducing the spatial resolution
- Like convolutional layer, *pooling* operators consist of a fixed-shape window that is slid over all regions of the input according to its stride
- However unlike convolutional layer, the pooling layer contains no parameters (there is no *kernel*)
- Mostly, we take the maximum of the elements in the pooling window – **max pooling**
- Max pooling gets feature representation that is somewhat invariant to translation (recall we wanted to find Waldo irrespective of its location in image)
- There is also **average pooling**, taking average of the elements in the window


- Now we have all the ingredients to assemble a CNN
- We will now see the first CNN — *LeNet* (1998) by Yann LeCun for handwritten digit recognition

# LeNet for handwritten digit recognition

- We have a **grayscale image** of an handwritten digit of size 32 x 32 with depth = 1
- This is going to be our input to LeNet



# LeNet for handwritten digit recognition




$$\left\lfloor \frac{32 - 5 - 0}{1} + 1 \right\rfloor = 28$$

- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 5 \times 5$
- # kernels  $\rightarrow 6$
- Parameters  $\rightarrow$

- We have 6 kernels
- Each kernel has  $5 \times 5 = 25$  weights
- # parameters  $= 25 \times 6 = 150$
- Input size  $= 32 \times 32 = 1024$
- Output size  $= 28 \times 28 = 784$
- If this was a fully-connected network, you needed  $1024 \times 784$  weights!
- For convolution layer, we have just 150 parameters
- Great reduction in # of parameters
- A sigmoid activation was applied (ReLU was not known then)

# LeNet for handwritten digit recognition



- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 1 \times 5 \times 5$
- # kernels  $\rightarrow 6$

• Parameters  $\rightarrow$

- Stride  $S = 2$
- Pad  $P = 0$
- Kernel  $\rightarrow 2 \times 2$

• Parameters  $\rightarrow$


- Max pooling is a per feature map operation
- Here the kernel size is  $2 \times 2$
- It downscales the size of feature maps

$$f_h = \left\lfloor \frac{28 - 2 + 0}{2} + 1 \right\rfloor = 14$$

$$f_w = \left\lfloor \frac{28 - 2 + 0}{2} + 1 \right\rfloor = 14$$

- The depth of max pooling layer is the same as the preceding convolution layer; here depth is 6
- There are no parameters in max pooling layers; they just take maximum of elements in a window

# LeNet for handwritten digit recognition



- Feature map size

$$f_h = \left\lfloor \frac{14 - 5 + 0}{1} + 1 \right\rfloor = 10$$

$$f_w = \left\lfloor \frac{14 - 5 + 0}{1} + 1 \right\rfloor = 10$$

- Depth of kernels = 6
- Here kernel size is 5 x 5 x 6
- # parameters = 5 x 5 x 6 x 16 = 2400

- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 5 \times 5$
- # kernels  $\rightarrow 6$

- Parameters  $\rightarrow$


- Stride  $S = 2$
- Pad  $P = 0$
- Kernel  $\rightarrow 2 \times 2$

- Parameters  $\rightarrow$

- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 5 \times 5 \times 6$
- # kernels  $\rightarrow 16$

- Parameters  $\rightarrow$

# LeNet for handwritten digit recognition



$$f_h = \left\lfloor \frac{10 - 2 + 0}{2} + 1 \right\rfloor = 5$$

$$f_w = \left\lfloor \frac{10 - 2 + 0}{2} + 1 \right\rfloor = 5$$

- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 5 \times 5$
- # kernels  $\rightarrow 6$

- Parameters →

- Stride  $S = 2$
- Pad  $P = 0$
- Kernel  $\rightarrow 2 \times 2$

- Parameters →

- Stride  $S = 1$
- Pad  $P = 0$
- Kernel  $\rightarrow 5 \times 5 \times 6$
- # kernels  $\rightarrow 16$


- Parameters →

- Stride  $S = 2$
- Pad  $P = 0$
- Kernel  $\rightarrow 2 \times 2$

- Parameters  $\rightarrow 0$



# LeNet for handwritten digit recognition



After max pooling layer 2, there are two fully connected hidden layers

- The features of the max pooling layer is flattened out into a vector of size  $16 \times 5 \times 5 = 400$  and fed to FC 1 layer as inputs
- FC 1 layer has 120 hidden units  $\rightarrow (16 \times 5 \times 5) \times 120 = 48000$  weights + 120 biases = 48120 parameters
- FC 2 layer has 84 hidden units  $\rightarrow 120 \times 84 = 10080$  weights + 84 biases = 10164 parameters
- Output layer has 10 classes  $\rightarrow 84 \times 10 = 840$  weights + 10 biases = 850 parameters
- The entire network can be trained using back propagation