

Lecture 17: Kernel Theory

With kernel ridge regression (KRR) and support vector regression (SVR) we learned three concepts:

1) Primal and dual formulations of a model

- Primal formulation expresses the model in terms of $\underline{\theta} \in \mathbb{R}^d$
- Dual formulation uses $\underline{\alpha} \in \mathbb{R}^N$ ($N \leftarrow$ size of training dataset), and does not depend on the value of 'd'
- Both formulations are mathematically equivalent
 - Primal formulation is useful if $N > d$
 - Dual formulation is useful if $d > N$

2> We introduced **kernels** $K(\underline{x}, \underline{x}')$ that allows us to let $d \rightarrow \infty$ without explicitly formulating an infinite vector of non-linear transformations $\underline{\phi}(\underline{x})$


- The dual formulation is particularly useful when using kernel methods, since the dimension of \underline{Q} in the primal formulation could be very large

3> We used different loss functions (and included L_2 -regularization)

- KRR makes use of squared error loss

- SVR uses ϵ -insensitive loss

→ gives sparse $\underline{\alpha}$
in the dual formulation



Kernel theory

Lets look a bit more into kernels

- Kernel was defined as being any function that takes in two arguments and returns a scalar
- We also suggested that we will restrict ourselves to ^{positive semi-definite} PSD kernels
- Vanilla kNN \rightarrow kernel kNN (provides a variety of distance metrics)
 - Recall that vanilla kNN constructs prediction for \underline{x}^* by taking the average or a majority vote among the k "nearest" neighbours
 - In its standard form, "nearest" was defined by the Euclidean distance
 - Euclidean distance between 2 points \underline{x} and \underline{x}' : $\|\underline{x} - \underline{x}'\|_2$ (always +ve)

Euclidean distance between 2 points \underline{x} and \underline{x}' : $\|\underline{x} - \underline{x}'\|_2$ (always +ve)

- Since Euclidean distance is positive, we can consider squared Euclidean distance instead

$$\begin{aligned}\|\underline{x} - \underline{x}'\|_2^2 &= (\underline{x} - \underline{x}')^T (\underline{x} - \underline{x}') \\ &= \underline{x}^T \underline{x} + \underline{x}'^T \underline{x}' - 2 \underline{x}^T \underline{x}'\end{aligned}$$

For many kernels, these terms are mostly constants (e.g. RBF kernel)

Define a kernel $\kappa(\underline{x}, \underline{x}') = \underline{x}^T \underline{x}'$

$$= \underline{\kappa(\underline{x}, \underline{x})} + \underline{\kappa(\underline{x}', \underline{x}')} - \underbrace{2 \kappa(\underline{x}, \underline{x}')}_{\text{this term determines how close any two points are}}$$

this term is more interesting

} $\kappa(\underline{x}, \underline{x}')$ takes a large value if \underline{x} and \underline{x}' are close

- In kernel kNN, $\kappa(\underline{x}, \underline{x}')$ can be replaced with any PSD kernel

- How can you use vanilla kNN where Euclidean distance has no natural meaning?

Example: Distance between words which reflect sentiment

Word	Sentiment
Tremendous	Positive
Horrific	Negative
Outrageous	Negative

- what could be the label for "horrendous"?

- One may think of converting the input space to numbers first and then use Euclidean distance

$x^* = \text{Horrendous}$

$k = 1 \rightarrow \text{Positive}$

$k = 3 \rightarrow \text{Negative}$

- An easier way to compare is using, for ex, Levenshtein distance (LD), which is the number of single-character edits needed to transform one word (string) into another

- One can construct a kernel as $K(x, x') = \exp\left(-\frac{(\text{LD}(x, x'))^2}{2l^2}\right)$

to implement kernel kNN (instead of vanilla kNN)

Lessons learned about kernels so far

- A kernel defines how close/similar any two points are
 - If $\kappa(\underline{x}^{(i)}, \underline{x}^*) > \kappa(\underline{x}^{(j)}, \underline{x}^*)$, then \underline{x}^* is more similar to $\underline{x}^{(i)}$ than $\underline{x}^{(j)}$
 - It also implies that prediction $\hat{y}(\underline{x}^*)$ is most influenced by the training data points that are closest to \underline{x}^*
 - Therefore, a kernel plays an important role of determining the individual influence of each training data point when making a prediction
- No need to bother about the inner product $\underline{\phi}(\underline{x})^T \underline{\phi}(\underline{x}')$ once we have introduced the kernel $\kappa(\underline{x}, \underline{x}')$

Lessons learned about kernels so far

- Choice of a kernel corresponds to preference for certain types of functions

– For example, the squared exponential (or RBF) kernel

$$K(\underline{x}, \underline{x}') = \exp\left(-\frac{\|\underline{x} - \underline{x}'\|_2^2}{2l^2}\right)$$

implies a preference for smooth functions

– In primal formulation, we choose features $\underline{\phi}(\underline{x})$ which will reflect the type of transformations we want to introduce. This choice is reflected to some extent in choosing kernels in the dual formulation

A machine learning engineer must choose a kernel wisely and should not simply resort to 'default' choices

What are valid choices of kernels?

- We already know that kernels are a way to represent non-linear feature transformation $\underline{\phi}(\underline{x})$

$$K(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x})^T \underline{\phi}(\underline{x}')$$

- Question: Does an arbitrary kernel $K(\underline{x}, \underline{x}')$ always correspond to a feature transformation $\underline{\phi}(\underline{x})$?
 - The question is primarily of theoretical nature
 - Practically, it matters very less whether a kernel $K(\underline{x}, \underline{x}')$ admits a factorization $K(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x})^T \underline{\phi}(\underline{x}')$ or not
 - Furthermore, the factorization has no direct correspondence to how well the kernel will perform in terms of E_{new} , which still has to be evaluated using cross-validation

Question: Does an arbitrary kernel $\kappa(\underline{x}, \underline{x}')$ always correspond to a feature transformation $\underline{\phi}(\underline{x})$?

Answer: Yes, if the kernel $\kappa(\underline{x}, \underline{x}')$ is PSD (positive semi-definite)
(no negative eigen-values)

Recall that a kernel is PSD if the Gram matrix $\underline{K}(\underline{X}, \underline{X})$ is PSD
for any \underline{X}

- It holds that any kernel $\kappa(\underline{x}, \underline{x}')$ that is defined as an inner product between feature vectors $\underline{\phi}(\underline{x})$ is always PSD

$$\begin{aligned}\kappa(\underline{x}, \underline{x}') &= \underline{\phi}(\underline{x})^T \underline{\phi}(\underline{x}') \\ &= \langle \underline{\phi}(\underline{x}), \underline{\phi}(\underline{x}') \rangle\end{aligned}$$

$\langle \cdot, \cdot \rangle \leftarrow$ inner product

Show $\underline{v}^T \underline{K}(\underline{X}, \underline{X}) \underline{v} \geq 0$ for any vector \underline{v} (do yourself)

Question: Does an arbitrary kernel $\kappa(\underline{x}, \underline{x}')$ always correspond to a feature transformation $\underline{\phi}(\underline{x})$?

Answer: Yes, if the kernel $\kappa(\underline{x}, \underline{x}')$ is PSD (positive semi-definite)
(no negative eigen-values)


- It holds that any kernel $\kappa(\underline{x}, \underline{x}')$ that is defined as an inner product between feature vectors $\underline{\phi}(\underline{x})$ is always PSD

$$\underbrace{\underline{\phi}(\underline{x})}_{\text{feature vector}} \xrightarrow{\text{inner product}} \underbrace{\kappa(\underline{x}, \underline{x}')}_{\text{PSD}}$$

- The other direction also holds true, that is, for any PSD kernel $\kappa(\underline{x}, \underline{x}')$ there always exist a feature vector $\underline{\phi}(\underline{x})$ such that $\kappa(\underline{x}, \underline{x}')$ can be written as its inner product

$$\underbrace{\underline{\phi}(\underline{x})}_{\text{feature vector}} \longleftarrow \underbrace{\kappa(\underline{x}, \underline{x}')}_{\text{if PSD}}$$

- The other direction also holds true, that is, for any PSD kernel $\kappa(\underline{x}, \underline{x}')$ there always exist a feature vector $\underline{\phi}(\underline{x})$ such that $\kappa(\underline{x}, \underline{x}')$ can be written as its inner product



- It can be shown that for any PSD kernel, it is possible to construct a function space, more specifically a Hilbert space, that is spanned by a feature vector $\underline{\phi}(\underline{x})$ s. t. $\kappa(\underline{x}, \underline{x}') = \underline{\phi}(\underline{x})^T \underline{\phi}(\underline{x}')$
 - There are multiple ways to construct a Hilbert space space spanned by $\underline{\phi}(\underline{x})$. One of the ways is using the so-called reproducing kernel Hilbert space (RKHS) mapping

A brief introduction to Reproducing Kernel Hilbert Spaces (RKHS)

- Euclidean space is a space of vectors equipped with inner products between vectors
- **Hilbert space** ^{→ space of functions with inner product} is a generalization of Euclidean space to functions (which can be treated as infinite dimensional vectors). It allows inner product between functions
- A Hilbert space H is called the **RKHS** if there exists a kernel $k(\underline{x}, \underline{x}')$ with the **reproducing property** that

$$f(\underline{x}') = \langle f(\cdot), k(\cdot, \underline{x}') \rangle \quad \forall f \in H, \quad \forall \underline{x}'$$

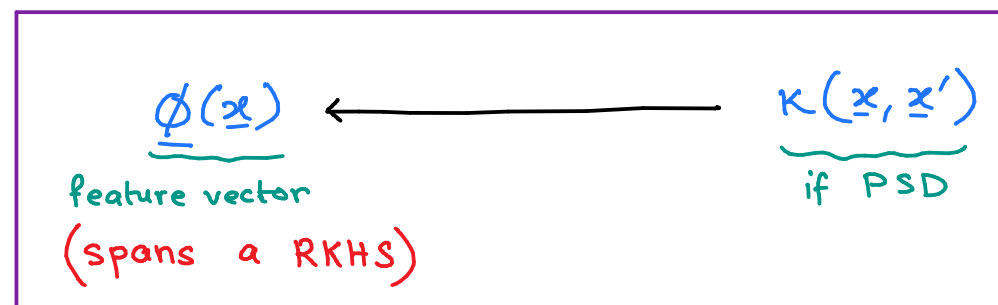
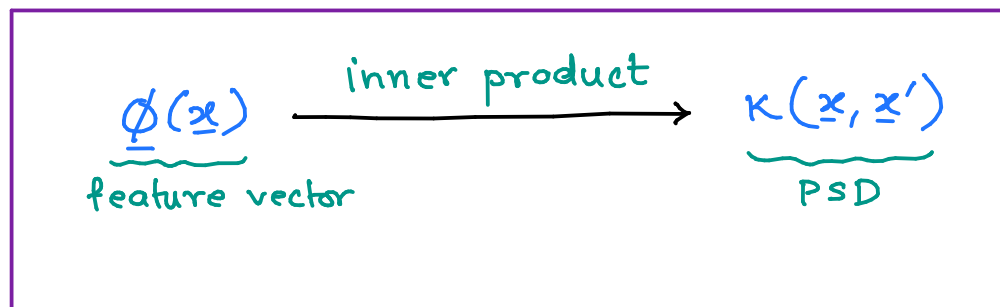
- If we set $f(\cdot) = k(\cdot, \underline{x})$, then

$$\langle k(\cdot, \underline{x}), k(\cdot, \underline{x}') \rangle = k(\underline{x}, \underline{x}')$$


This reproducing property is the main building block of RKHS. This RKHS is spanned by the corresponding feature $\phi(\underline{x})$ of kernel $k(\underline{x}, \underline{x}')$

Question: Does an arbitrary kernel $\kappa(\underline{x}, \underline{x}')$ always correspond to a feature transformation $\underline{\phi}(\underline{x})$?

Answer: Yes, if the kernel $\kappa(\underline{x}, \underline{x}')$ is PSD (positive semi-definite)
(no negative eigen-values)



- A given Hilbert space uniquely defines a kernel, but for a kernel there exists multiple Hilbert spaces which correspond to it



E.g. $\kappa(\underline{x}, \underline{x}') = \underline{x}^T \underline{x}'$

\searrow
 $\underline{\phi}_1(\underline{x}) = \underline{x}$
 (one-dimensional)

\searrow
 $\underline{\phi}_2(\underline{x}) = \begin{bmatrix} \underline{x}/\sqrt{2} \\ \underline{x}/\sqrt{2} \end{bmatrix}$
 (two-dimensional)