

Introduction to Ensemble methods

- We have looked at different machine learning models now
 - k-Nearest Neighbours
 - Classification & Regression Trees (CART)
 - Logistic Regression
 - Neural Networks
 - Support Vector Machines
- Question: Given a dataset, which ML algorithm should we pick, and how do you know which technique will perform the best?
- Unfortunately, there is no good answer to this question.
 - It is mostly a process of trial-and-error
 - Each kind of ML algorithm yields a different model/hypothesis
 - But there is no perfect model/hypothesis in practice
- So you may ask could we combine several imperfect models into a better model?

- Analogies of combining multiple models in our society
 - Elections combine voter's choices to pick a "good" candidate
 - Committees combine several experts' opinion to make better decisions
- Intuition behind combining multiple models/hypotheses
 - Individuals (or individual models) often make mistakes, but the "majority" is less likely to make mistakes
 - Individuals often have partial knowledge, but a committee can pool expertise to make better decisions
- Ensemble learning can combine an ensemble of
 - Different types of base models (e.g. Neural networks, CART and SVM)
 - Same base model trained slightly differently ✓ We are going to follow this approach

Bagging (or Bootstrap Aggregating)


- A central concept in ML is the bias-variance tradeoff
 - The more flexible a model is, the lower its bias will be
- Examples of highly flexible models that can represent complicated input-output relationships are k-Nearest Neighbours, CART, NNs, etc.
- The downside of such highly flexible models is the risk of overfitting
- Overfitted models lead to unwanted high variance in predictions
- By using bagging, we can reduce the variance of the base model without increasing its bias
- Lets take an example of regression trees with bagging

- Consider the data obtained as

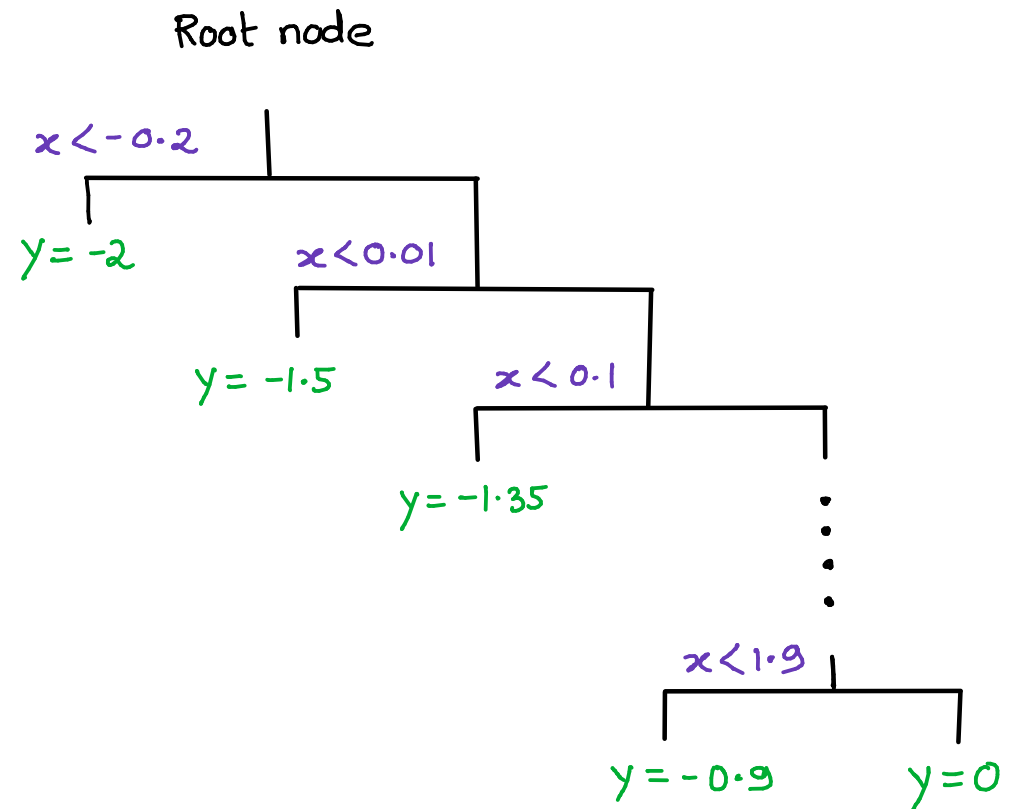

$$y = f(x) + \epsilon$$

↖ noise



- We would like to train an ML model using this data, so as to be able to predict new data points well
- A good prediction would mean that the trained model should predict $f(x)$ shown by the dotted line well at x^*
- For this problem, let us use Regression Trees as the chosen ML method, since they are non-parametric methods and are very flexible



- Recall that in Classification and Regression trees, we partition the input space using box-shaped decision boundaries
- Lets consider a Regression tree which is grown until each leaf node has only one data point




- On fitting a regression tree (with one data point in each leaf), we get an **OVERFITTED** Regression Tree



- Due to overfitting, the resulting regression tree is a **low-bias-high-variance** model
 - high variance means the trained model is very sensitive to the training data; if the training data changes, the predictions change a lot

- Because of the noise in training data, we can think of the prediction $\hat{y}(x^*)$ from the trained model as a **random variable**
 - It means that if we had multiple datasets and we trained different RTs on them, each of their predictions $\hat{y}(x^*)$ would be different



- So if we assumed that we had access to B **independent** datasets $T^{(1)}, T^{(2)}, \dots, T^{(B)}$, then we could train a separate tree for each dataset and obtain separate predictions $\hat{y}_b(x^*)$, $b=1, 2, \dots, B$, then:
 - Each $y_b(x^*)$ would have **low bias** and **high variance**
 - By averaging $\hat{y}(x^*) = \frac{1}{B} \sum_{b=1}^B \hat{y}_b(x^*)$, the bias is kept small, but the variance is reduced by a factor of B ! (Proof?)

Probability detour - Variance reduction by averaging

Let z_1, z_2, \dots, z_B be a collection of identically distributed but possibly dependent random variables, with

$$\left. \begin{array}{l} \text{Mean: } \mathbb{E}[z_b] = \mu \\ \text{Variance: } \text{Var}(z_b) = \sigma^2 \end{array} \right\} \text{ for } b = 1, 2, \dots, B$$
$$\text{Correlation: } \text{Corr}(z_i, z_j) = \rho \quad i \neq j, \quad i, j = 1, 2, \dots, B$$

Then one can show that the mean and variance of the average $\frac{1}{B} \sum_{b=1}^B z_b$ are: (assuming $\rho > 0$)

$$\mathbb{E} \left[\frac{1}{B} \sum_{b=1}^B z_b \right] = \mu, \quad \text{Var} \left[\frac{1}{B} \sum_{b=1}^B z_b \right] = \frac{1-\rho}{B} \sigma^2 + \rho \sigma^2$$

small for large B

- **Problem**: We only have access to **one training dataset**
 - **Solution**: **Bootstrap** the data!
 - Bootstrap is a method of artificially creating multiple datasets (of size N) out of one dataset (also of size N)
 - Sample N times with replacement from the original training data $\mathcal{T} = \{ \underline{x}^{(i)}, y^{(i)} \}_{i=1}^N$
 - Repeat B times to generate B "bootstrapped" training datasets $\tilde{\mathcal{T}}^{(1)}, \tilde{\mathcal{T}}^{(2)}, \dots, \tilde{\mathcal{T}}^{(B)}$
 - **BAGGING**
 - For each bootstrapped dataset $\tilde{\mathcal{T}}^{(b)}$, we train a tree (basemodel)
- Averaging them,

$$\hat{y}_{\text{bag}} = \frac{1}{B} \sum_{b=1}^B \tilde{y}^b(\underline{x})$$

Bagging example with regression trees as basemodel

Assume that we have a training set

$$\mathcal{T} = \left\{ (\underline{x}^{(1)}, y^{(1)}), (\underline{x}^{(2)}, y^{(2)}), (\underline{x}^{(3)}, y^{(3)}), \dots, (\underline{x}^{(N)}, y^{(N)}) \right\}$$

- We generate, say, $B = 9$ datasets by bootstrapping:

$$\tilde{\mathcal{T}}^{(1)} = \left\{ (\underline{x}^{(1)}, y^{(1)}), (\underline{x}^{(2)}, y^{(2)}), (\underline{x}^{(3)}, y^{(3)}), \dots, (\underline{x}^{(3)}, y^{(3)}) \right\}$$

$$\tilde{\mathcal{T}}^{(2)} = \left\{ (\underline{x}^{(1)}, y^{(1)}), (\underline{x}^{(N)}, y^{(N)}), (\underline{x}^{(N)}, y^{(N)}), \dots, (\underline{x}^{(N)}, y^{(N)}) \right\}$$

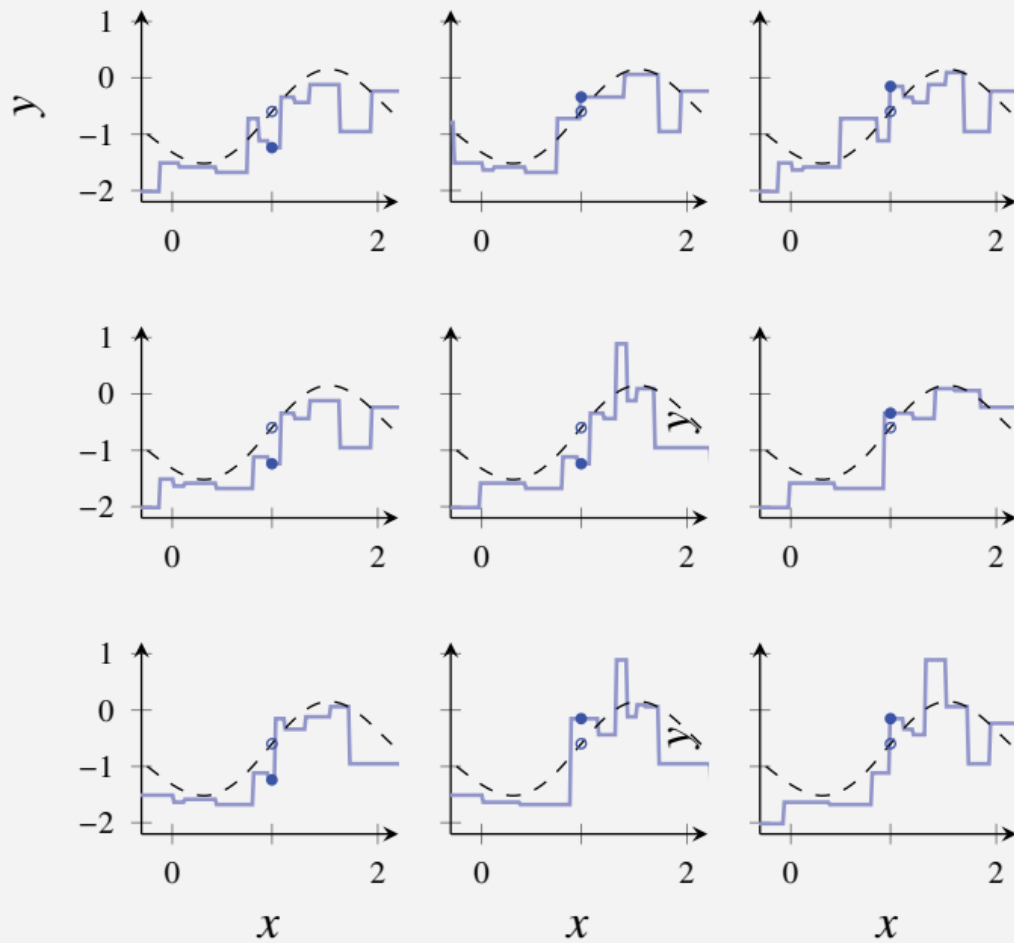
$$\vdots$$

$$\tilde{\mathcal{T}}^{(9)} = \left\{ (\underline{x}^{(1)}, y^{(1)}), (\underline{x}^{(1)}, y^{(1)}), (\underline{x}^{(2)}, y^{(2)}), \dots, (\underline{x}^{(3)}, y^{(3)}) \right\}$$


- We compute $B = 9$ (deep) regression trees $\tilde{y}^{(1)}(\underline{x}), \tilde{y}^{(2)}(\underline{x}), \dots, \tilde{y}^{(9)}(\underline{x})$
one for each dataset $\tilde{\mathcal{T}}^{(1)}, \tilde{\mathcal{T}}^{(2)}, \dots, \tilde{\mathcal{T}}^{(9)}$, and average $\hat{y}_{\text{bag}} = \frac{1}{9} \sum_{b=1}^9 \tilde{y}^{(b)}(\underline{x})$

Bagging example with regression trees


Ensemble of bootstrapped regression trees



Regression tree with complete data



Average of bootstrapped regression trees = bagging



Bagging algorithm

- Training : Learn all base models

Data: Training dataset $\mathcal{T} = \{ \underline{x}^{(i)}, y^{(i)} \}_{i=1}^N$

Result: 'B' base models

for $b = 1, \dots, B$ do

- Generate a bootstrap dataset $\tilde{\mathcal{T}}^{(b)} = \{ \tilde{\underline{x}}^{(i)}, \tilde{y}^{(i)} \}_{i=1}^N$

- Learn a base model from $\tilde{\mathcal{T}}^{(b)}$

end

Obtain $\hat{y}_{\text{bag}}(\underline{x})$ by averaging: $\hat{y}_{\text{bag}}(\underline{x}) = \frac{1}{B} \sum_{i=1}^B \tilde{y}^{(b)}(\underline{x})$

- Prediction with the base models

Data: 'B' base models and test input \underline{x}^*

Result: A prediction $\hat{y}_{\text{bag}}(\underline{x}^*)$

Use same formula

The choice of 'B' is mainly guided by computational constraints

RANDOM FORESTS

- Bagging can greatly improve the performance of CART
 - Averaging over ensemble prediction, in case of regression trees
 - Majority vote over ensemble prediction, for classification trees
- However, the 'B' bootstrapped dataset are **correlated**!

Therefore, the variance reduction due to averaging is diminished

Recall

$$\text{Var} \left[\frac{1}{B} \sum_{b=1}^B z_b \right] = \frac{1-\rho}{B} \sigma^2 + \rho \sigma^2$$

– No variance reduction when $\rho = 1$

– Highest variance reduction when $\rho = 0$

- **Idea of Random Forest**: De-correlate the 'B' trees by injecting additional randomness when constructing each tree

Random forest algorithm

Inputs: $\mathcal{T} = \{ \mathbf{x}^{(i)}, y^{(i)} \}_{i=1}^N$; $\mathbf{x} \in \mathbb{R}^p$

for $b=1$ to B , do (can run in parallel)

(a) Draw a bootstrap dataset $\tilde{\mathcal{T}}^{(b)}$ of size N from \mathcal{T}

(b) Grow a regression (or classification) tree by repeating the steps below, until a minimum node size is reached:

- Select a random subset consisting of $q \leq p$ inputs
- Find the best splitting variable x_j among the 'q' selected inputs
- Split the node into two children with $\{x_j \leq s\}$ and $\{x_j > s\}$

Thumb rule
 $q = \sqrt{p}$ (for CT)
 $q = p/3$ (for RT)

Final model is the average of the 'B' ensemble members

$$\hat{y}^{rf} = \frac{1}{B} \sum_{b=1}^B \tilde{y}^{(b)}$$

- For identically distributed random variables $\{z_b\}_{b=1}^B$

$$\text{Var} \left[\frac{1}{B} \sum_{b=1}^B z_b \right] = \frac{1-\rho}{B} \sigma^2 + \rho \sigma^2$$

- The random input selection used in random forests:
 - increases the bias, but often very slowly ↓
 - adds to the variance (σ^2) of each tree ↓
 - reduces the correlation (ρ) between member trees ↑↑↑
- The reduction in correlation typically has a dominant effect
⇒ leads to an overall reduction in error
- Bagging is a general technique → can be used with any base model
Random forest consider base models as classification or regression trees